

Discovery at Scale: Understanding an Existing Hierarchical Codebase


Before you write a single spec, you have to understand what's already there.

What You'll Learn

- 1 Explain why reading code alone doesn't scale on a large, hierarchical codebase
- 2 Use MCP to connect an agent to design documentation and a bug/defect tracker
- 3 Scope discovery to a service boundary and produce a lightweight dependency map

THE CHALLENGE


Why Code Alone Doesn't Scale



Gray = one of 50+ interconnected services · Coral = the file you can read right now

THE APPROACH


MCP-Connected Discovery



MCP connects the agent to what a code search alone would never find

SCOPE IT

Scope discovery to one service boundary, not the whole repository.



Everything outside the boundary stays out of scope for this pass

APPLYING THE CONCEPTS

GitHub Spec Kit

No Dedicated Discovery Command

No direct equivalent

Spec Kit's command chain starts at constitution and specify. It has no built-in command for reading unfamiliar code or reconciling it against external documentation before a spec exists, because Spec Kit assumes you already understand what you're specifying.

What to use instead

Discovery happens through your coding agent's own MCP connections (Copilot or Claude Code), one layer below the Spec Kit CLI. The resulting context is what you hand to `/speckit.specify` as background when you write the spec in Module 3, so nothing here is wasted, it just isn't a named Spec Kit command.

APPLYING THE CONCEPTS

OpenSpec

Explore: Built for This Moment

```
/opsx:explore
```

A no-stakes thinking partner: reads your codebase, compares options, and sharpens a fuzzy idea into a concrete plan before any change exists, so you can test a direction without committing to it yet.

With MCP: With MCP connections in place, explore can incorporate design docs and tracker context directly into its investigation, which is what makes it a genuine fit for this module's discovery work.

No artifacts yet: Explore creates no artifacts by itself: it hands off to /opsx:propose once the picture is clear, keeping the exploratory, low-commitment phase clearly separate from the moment you start proposing a real change.

Module 2 at a Glance

	GitHub Spec Kit	OpenSpec
Dedicated command	None (handled by the agent's MCP layer)	/opsx:explore
Output	Context fed into the next spec	Investigation only, no artifact
Fit for this course	Discovery happens upstream of the CLI	Purpose-built for brownfield investigation

Key Takeaways

- Discovery precedes specification, always, on a brownfield codebase, because you can't write requirements against a system you don't yet understand.
- MCP is the connective tissue between your agent and everything outside the repo: docs, trackers, and prior decisions.
- OpenSpec's explore command is a natural fit for this step; Spec Kit expects discovery to happen upstream of it, in your agent rather than in the CLI.