

DAY 3

Module 7

Multi-Agent Workflows and Team Adoption

Coordinating parallel work, and keeping an audit trail across teams and toolchains.

What You'll Learn

- 1 Coordinate multiple agents or subagents across a task graph
- 2 Decide when to parallelize versus sequence work
- 3 Apply repository conventions and governance across teams and toolchains

Parallelize or Sequence?

Some work parallelizes cleanly: independent services, independent tasks. Some doesn't: shared state, sequential dependencies.

Delegating a focused sub-task to a second agent instance works best when the boundary is clear and the output is easy to review on its own.

Conventions and Governance

A specs/ subdirectory per service, with specs linked to issues and pull requests.

Reviewing agent actions before execution, managing secrets and permissions, and keeping an audit trail that holds up across teams using different toolchains.

WHY IT MATTERS

Uncoordinated parallel agents create their own kind of chaos.

Two agents editing overlapping code without visibility into each other produce conflicts no single review catches early.

An audit trail only works if every agent's actions were reviewable in the first place.

APPLYING THIS WITH

GitHub Spec Kit

Workflows: Explicit Orchestration

`specify workflow run`

YAML-defined workflows chain commands, prompts, and shell steps into repeatable sequences, with gate steps for human checkpoints.

Parallel primitive: fan-out and fan-in step types dispatch a step across a list of items and aggregate the results, a direct primitive for coordinating parallel agent work across a task graph.

Audit trail: Every run persists state (state.json, log.jsonl) and can pause and resume from the exact point of interruption, doubling as an audit trail.

APPLYING THIS WITH

OpenSpec

A partial fit: some of this module's concept maps, some doesn't.

Bulk-Archive, Not Orchestration

```
/opsx:bulk-archive
```

Coordinates multiple parallel change streams at once, detecting and resolving spec conflicts between them by checking what's actually implemented.

Worth knowing: OpenSpec's beta "Stores" feature extends planning across multiple repos or teams, for organizations working past a single-repo scope.

Important distinction: Genuine agent-to-agent delegation is a capability of the underlying coding agent (Claude Code subagents, Copilot Fleet mode), not a feature either SDD framework owns directly.

Module 7 at a Glance

	GitHub Spec Kit	OpenSpec
Orchestration primitives	fan-out / fan-in, gates, resumable state	None (conflict-aware bulk archiving only)
Multi-stream support	Workflow YAML with conditional logic	/opsx:bulk-archive
Agent delegation	Not framework-owned (agent capability)	Not framework-owned (agent capability)

Key Takeaways

- Parallelism is a property of the work, not a default setting.
- Spec Kit's workflow engine gives you explicit orchestration primitives; OpenSpec's support is lighter-weight.
- Delegation to a second agent is your coding agent's job, not the SDD framework's.

HANDS-ON NEXT

Propose and review an agent-generated plan for your project's next multi-step increment before execution. Delegate a focused sub-task to a second agent instance, then critique the delegated output and document what review step should have caught any issues.

Next: Module 8: Common Pitfalls and Anti-Patterns